

STTAI - Lab Assignment 10

| Name | Roll Number |
|--------------------|-------------|
| Romit Mohane | 23110279 |
| Rudra Pratap Singh | 23110281 |

Using real-world data, this assignment will introduces us to key concepts in A/B testing and Covariate Shift Detection. We performed hypothesis testing using the scipy library and identified distributional shifts in datasets using classification-based techniques.

Part 1: A/B Testing using Ad Click Prediction

1. Load the Dataset into a pandas df

```
df = pd.read_csv('ad_click_dataset.csv')

df
```

| | id | full_name | age | gender | device_type | ad_position | browsing_history | time_of_day | click |
|------|------|-----------|------|------------|-------------|-------------|------------------|-------------|-------|
| 0 | 670 | User670 | 22.0 | NaN | Desktop | Top | Shopping | Afternoon | 1 |
| 1 | 3044 | User3044 | NaN | Male | Desktop | Top | NaN | NaN | 1 |
| 2 | 5912 | User5912 | 41.0 | Non-Binary | NaN | Side | Education | Night | 1 |
| 3 | 5418 | User5418 | 34.0 | Male | NaN | NaN | Entertainment | Evening | 1 |
| 4 | 9452 | User9452 | 39.0 | Non-Binary | NaN | NaN | Social Media | Morning | 0 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 9995 | 8510 | User8510 | NaN | NaN | Mobile | Top | Education | NaN | 0 |
| 9996 | 7843 | User7843 | NaN | Female | Desktop | Bottom | Entertainment | NaN | 0 |
| 9997 | 3914 | User3914 | NaN | Male | Mobile | Side | NaN | Morning | 0 |
| 9998 | 7924 | User7924 | NaN | NaN | Desktop | NaN | Shopping | Morning | 1 |
| 9999 | 3056 | User3056 | 44.0 | Male | Tablet | Top | Social Media | Morning | 0 |

10000 rows × 9 columns

2. Perform necessary data cleaning and preprocessing:

df_nullhandled_dropped

| | id | full_name | age | gender | device_type | ad_position | browsing_history | time_of_day | click |
|------|------|-----------|------|------------|-------------|-------------|------------------|-------------|-------|
| 17 | 188 | User188 | 56.0 | Female | Tablet | Bottom | News | Morning | 1 |
| 25 | 4890 | User4890 | 43.0 | Male | Tablet | Bottom | Education | Afternoon | 1 |
| 33 | 4985 | User4985 | 37.0 | Male | Mobile | Top | News | Evening | 0 |
| 52 | 9888 | User9888 | 49.0 | Male | Mobile | Top | News | Morning | 1 |
| 102 | 8201 | User8201 | 59.0 | Female | Desktop | Bottom | Social Media | Morning | 0 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 9951 | 7268 | User7268 | 28.0 | Female | Desktop | Bottom | News | Evening | 1 |
| 9952 | 5912 | User5912 | 41.0 | Non-Binary | Mobile | Side | Education | Night | 1 |
| 9960 | 9638 | User9638 | 64.0 | Non-Binary | Desktop | Top | Entertainment | Morning | 0 |
| 9986 | 5574 | User5574 | 52.0 | Female | Desktop | Bottom | Shopping | Afternoon | 1 |
| 9999 | 3056 | User3056 | 44.0 | Male | Tablet | Top | Social Media | Morning | 0 |

816 rows × 9 columns

```
# Work on the dropped values DF
df_current = df_nullhandled_dropped.copy()
# Keep only the two positions of interest
df_current = df_current[df_current['ad_position'].isin(['Top','Bottom'])]

# Map to numbers
df_current['ad_position_flag'] = df_current['ad_position'].map({'Top': 0, 'Bottom': 1})
df_current['gender'] = df_current['gender'].map({'Male': 0, 'Female': 1, 'Non-Binary': 2})
df_current['time_of_day'] = df_current['time_of_day'].map({'Morning': 0, 'Afternoon': 1, 'Evening': 2, 'Night': 3})
df_current['device_type'] = df_current['device_type'].map({'Desktop': 0, 'Mobile': 1, 'Tablet': 2})
df_current['browsing_history'] = df_current['browsing_history'].map({'Shopping': 0, 'Social Media': 1, 'Education': 2, 'News': 3, 'Entertainment': 4})
```

df_current # df with dropped rows where null values were present

| | id | full_name | age | gender | device_type | ad_position | browsing_history | time_of_day | click | ad_position_flag |
|------|------|-----------|------|--------|-------------|-------------|------------------|-------------|-------|------------------|
| 17 | 188 | User188 | 56.0 | 1 | 2 | Bottom | 3 | 0 | 1 | 1 |
| 25 | 4890 | User4890 | 43.0 | 0 | 2 | Bottom | 2 | 1 | 1 | 1 |
| 33 | 4985 | User4985 | 37.0 | 0 | 1 | Top | 3 | 2 | 0 | 0 |
| 52 | 9888 | User9888 | 49.0 | 0 | 1 | Top | 3 | 0 | 1 | 0 |
| 102 | 8201 | User8201 | 59.0 | 1 | 0 | Bottom | 1 | 0 | 0 | 1 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 9928 | 7790 | User7790 | 43.0 | 2 | 1 | Top | 1 | 0 | 0 | 0 |
| 9951 | 7268 | User7268 | 28.0 | 1 | 0 | Bottom | 3 | 2 | 1 | 1 |
| 9960 | 9638 | User9638 | 64.0 | 2 | 0 | Top | 4 | 0 | 0 | 0 |
| 9986 | 5574 | User5574 | 52.0 | 1 | 0 | Bottom | 0 | 1 | 1 | 1 |
| 9999 | 3056 | User3056 | 44.0 | 0 | 2 | Top | 1 | 0 | 0 | 0 |

558 rows × 10 columns

3. Split the dataset

```
groupA = df_current[df_current['ad_position_flag']==0]

# Group B: Bottom (flag=1)
groupB = df_current[df_current['ad_position_flag']==1]
```

groupA

| | id | full_name | age | gender | device_type | ad_position | browsing_history | time_of_day | click | ad_position_flag |
|------|------|-----------|------|------------|-------------|-------------|------------------|-------------|-------|------------------|
| 33 | 4985 | User4985 | 37.0 | Male | Mobile | Top | News | Evening | 0 | 0 |
| 52 | 9888 | User9888 | 49.0 | Male | Mobile | Top | News | Morning | 1 | 0 |
| 158 | 3007 | User3007 | 42.0 | Male | Desktop | Top | Shopping | Night | 0 | 0 |
| 204 | 8530 | User8530 | 52.0 | Female | Mobile | Top | Social Media | Afternoon | 1 | 0 |
| 231 | 4625 | User4625 | 33.0 | Non-Binary | Mobile | Top | News | Morning | 0 | 0 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 9888 | 5055 | User5055 | 45.0 | Male | Desktop | Top | Education | Morning | 1 | 0 |
| 9915 | 3335 | User3335 | 24.0 | Male | Mobile | Top | Entertainment | Night | 1 | 0 |
| 9928 | 7790 | User7790 | 43.0 | Non-Binary | Mobile | Top | Social Media | Morning | 0 | 0 |
| 9960 | 9638 | User9638 | 64.0 | Non-Binary | Desktop | Top | Entertainment | Morning | 0 | 0 |
| 9999 | 3056 | User3056 | 44.0 | Male | Tablet | Top | Social Media | Morning | 0 | 0 |

275 rows × 10 columns

groupB

| | id | full_name | age | gender | device_type | ad_position | browsing_history | time_of_day | click | ad_position_flag |
|-----|------|-----------|------|--------|-------------|-------------|------------------|-------------|-------|------------------|
| 17 | 188 | User188 | 56.0 | Female | Tablet | Bottom | News | Morning | 1 | 1 |
| 25 | 4890 | User4890 | 43.0 | Male | Tablet | Bottom | Education | Afternoon | 1 | 1 |
| 102 | 8201 | User8201 | 59.0 | Female | Desktop | Bottom | Social Media | Morning | 0 | 1 |

4. Use the statsmodel’s proportions_ztest function to perform an independent two-sample z-test between Group A and Group B.

```
clicks = [ groupA['click'].sum(),
           groupB['click'].sum() ]

# number of users in each group
nobs    = [ len(groupA),
           len(groupB) ]
print(clicks[0]/nobs[0], clicks[1]/nobs[1], "\n")

z_score, p_value = proportions_ztest(count=clicks, nobs=nobs)
print("z-score:", z_score)
print("p-value:", p_value)

if p_value < 0.05:
    print("Reject null hypothesis: the two CTRs are not equal.")
else:
    print("Fail to reject null hypothesis: the two CTRs are equal.")
```

0.6327272727272727 0.6784452296819788

z-score: -1.1365075404030447
p-value: 0.2557442115851094
Fail to reject null hypothesis: the two CTRs are equal.

5. **Interpret the result: Is there a statistically significant difference in click-through rates between the two groups? Justify your answer.** In our A/B test on 10,000 users (with missing values dropped), we compared click-through rates (CTRs) for ads shown at the Top vs. Bottom positions. Using a two-sample z-test, we obtained $z = -1.137$ and $p = 0.256$ (> 0.05), so we accept H_0 that the two CTRs are equal. The negative z-score indicates Bottom-positioned ads achieved a lower CTR than Top-positioned ads.

But, this difference is statistically insignificant.

Part 2: Covariate Shift Detection Using Air Quality Data

1. Load the datasets

```
import pandas as pd

air_df_train = pd.read_csv('train.csv')
air_df_test1 = pd.read_csv('test1.csv')
air_df_test2 = pd.read_csv('test2.csv')

air_df_train = air_df_train[air_df_train['NO2(GT)'] >= 0]
air_df_test1 = air_df_test1[air_df_test1['NO2(GT)'] >= 0]
air_df_test2 = air_df_test2[air_df_test2['NO2(GT)'] >= 0]

air_df_train.head()
```

| Unnamed: 0 | Date | Time | CO(GT) | PT08.S1(CO) | NMHC(GT) | C6H6(GT) | PT08.S2(NMHC) | NOx(GT) | PT08.S3(NOx) | NO2(GT) | PT08.S4(NO2) | PT08.S5(O3) | T | RH | AH | Unnamed: 15 | Unnamed: 16 | |
|------------|------|------------|----------|-------------|----------|----------|---------------|---------|--------------|---------|--------------|-------------|--------|------|------|-------------|-------------|-----|
| 1 | 2533 | 24/06/2004 | 07.00.00 | 1.2 | 1030.0 | -200.0 | 6.9 | 851.0 | 102.0 | 824.0 | 68.0 | 1700.0 | 983.0 | 21.9 | 57.0 | 1.4742 | NaN | NaN |
| 2 | 3047 | 15/07/2004 | 17.00.00 | 3.2 | 1164.0 | -200.0 | 20.3 | 1306.0 | 259.0 | 648.0 | 198.0 | 1886.0 | 1218.0 | 35.5 | 19.1 | 1.0888 | NaN | NaN |
| 3 | 805 | 13/04/2004 | 07.00.00 | 3.9 | 1496.0 | 524.0 | 19.1 | 1272.0 | 328.0 | 667.0 | 130.0 | 2011.0 | 1399.0 | 11.0 | 64.2 | 0.8398 | NaN | NaN |
| 4 | 2962 | 12/07/2004 | 04.00.00 | -200 | 780.0 | -200.0 | 1.8 | 568.0 | 24.0 | 1200.0 | 34.0 | 1331.0 | 501.0 | 19.9 | 51.3 | 1.1803 | NaN | NaN |
| 7 | 1910 | 29/05/2004 | 08.00.00 | 0.6 | 843.0 | -200.0 | 4.0 | 712.0 | 77.0 | 1303.0 | 77.0 | 1343.0 | 548.0 | 18.9 | 42.1 | 0.9081 | NaN | NaN |

```
air_df_train.columns
```

```
Index(['Unnamed: 0', 'Date', 'Time', 'CO(GT)', 'PT08.S1(CO)', 'NMHC(GT)',
      'C6H6(GT)', 'PT08.S2(NMHC)', 'NOx(GT)', 'PT08.S3(NOx)', 'NO2(GT)',
      'PT08.S4(NO2)', 'PT08.S5(O3)', 'T', 'RH', 'AH', 'Unnamed: 15',
      'Unnamed: 16'],
      dtype='object')
```

2. For each test dataset (test1.csv and test2.csv), compare it with train.csv using the Kolmogorov–Smirnov test (scipy.stats.ks_2samp). Perform the KS test on the NO2(GT) column to identify whether there are any distributional differences

```
ks_statistic_test1, p_value_test1 = ks_2samp(air_df_train['NO2(GT)'], air_df_test1['NO2(GT)'])
print("\nMean for test set 1:", air_df_test1['NO2(GT)'].mean())
print("KS Test for test1.csv:")
print("KS Statistic: {ks_statistic_test1}")
print("P-value: {p_value_test1}")

# Perform KS test for test2.csv
ks_statistic_test2, p_value_test2 = ks_2samp(air_df_train['NO2(GT)'], air_df_test2['NO2(GT)'])
print("\nMean for test set 2:", air_df_test2['NO2(GT)'].mean())
print("KS Test for test2.csv:")
print("KS Statistic: {ks_statistic_test2}")
print("P-value: {p_value_test2}")

if p_value_test1 < 0.05:
    print("\nReject the null hypothesis for test1.csv")
else:
    print("\nFail to reject the null hypothesis for test1.csv")

if p_value_test2 < 0.05:
    print("Reject the null hypothesis for test2.csv")
else:
    print("Fail to reject the null hypothesis for test2.csv")
```

```
Mean for train set: 94.57946026986507

Mean for test set 1: 94.53262518968134
KS Test for test1.csv:
KS Statistic: 0.017062220028073977
P-value: 0.9971378232852736

Mean for test set 2: 134.7030456852792
KS Test for test2.csv:
KS Statistic: 0.3688536442438679
P-value: 2.53172387531317e-74

Fail to reject the null hypothesis for test1.csv
Reject the null hypothesis for test2.csv
```

Therefore, there is a distributional difference in the values of **NO2(GT)** between the test sets.

3. Determine which of the two test datasets (test1.csv or test2.csv) exhibits a covariate shift relative to the training dataset (train.csv). Use the results of the Kolmogorov–Smirnov test to support your answer. The 2nd test set exhibits a covariate shift relative to the training set, since:

- The p-value for test set 1 and train set is **0.99714**
- The p-value for test set 2 and train set is around **0**

This rejects the Null Hypothesis for Test2 and shows strong covariate shift in **test2** dataset with respect to the **train** set.